DS8001 Project Proposal: Permutation-Based Approximate kNN Search for Mixed-Type Hospital Encounters

Jared Smouter and Amarnadh Nalla | Project path: **Empirical Analysis** | Group of two

1. Problem and questions. Exact k-nearest-neighbor (kNN) search computes a distance to all N indexed encounters. We study, as a purely algorithmic question (no clinical claim), how permutation-based approximate indexes trade recall against distance computations, index size and build cost under a mixed-type distance. The dataset is widely used for readmission prediction; our contribution is an algorithmic study of approximate similarity search on it. **RQ1.** How do the number of pivots n, the retained permutation length L and the re-ranking budget c affect recall@10 and query cost, and how do these costs scale with N? **RQ2.** At matched L and c, when does a truncated MI-File beat a truncated permutation scan, and how does random versus farthest-first pivot selection change this?

2. Methods. Implemented by us: (i) **brute-force exact kNN** (ground truth, ties broken by encounter id); (ii) a **truncated permutation scan** (TPS): each object stores the ranks of its L nearest pivots (rank L+1 otherwise); the query scores all objects by the Spearman footrule over the n pivots, keeps the best c and re-ranks them with the true distance (exact only within the candidate set); (iii) a **truncated MI-File** (Amato and Savino, 2008): per-pivot posting lists of (object, rank), with the same score and re-ranking, so only objects sharing a retained pivot are touched; (iv) random versus farthest-first pivot selection. Library baseline (not counted as ours): scikit-learn **BallTree** with the Manhattan metric on an exact weighted-L1 encoding of the distance, validated against brute force. Complexity analysis: build $O(Nn)$ distances plus ranking, storage $O(NL)$, query n pivot distances plus candidate scoring plus c re-ranking distances. **Distance:** six equal-weight feature groups (age, utilization, admission/discharge, ICD-9 diagnosis chapters, labs, medications); numeric features use scale-normalized absolute difference, nominal ones 0/1 mismatch, missing values get a fixed state or median plus indicator (no pairwise omission). With fixed scales and weights this is a metric, which we verify (identity, triangle inequality) on sampled triples.

3. Data. Diabetes 130-US Hospitals (1999–2008) (Strack et al., 2014; UCI, CC BY 4.0): 101,766 encounters, 71,518 patients. We exclude IDs, `readmitted`, race, gender and `weight`. Patients are split 90/5/5 into index/validation/test with a fixed seed, so no patient appears in both index and queries; scales and medians are fit on the index partition only.

4. Experiments and metrics. Nested index sizes N = 10k, 50k and the full index partition; 300 validation and 300 test queries (one encounter per patient). Grid: $n \in \{32, 64\}$, $L \in \{8, 16, 32\}$ ($L \le n$), $c/N \in \{0.5\%, 1\%, 5\%\}$, random and farthest-first pivots, 3 seeds each. We choose Pareto-optimal settings on validation, freeze them, and evaluate on test once. Metrics: recall@10 against brute force; distance computations, entries scored and postings read per query (machine-independent); serialized index size; build distances; median latency and peak memory (single-threaded, same machine). We report patient-level bootstrap 95% CIs, paired per-query differences, and seed variability separately from query variability. Deliverables: a Python 3.11+ package with README and fixed seeds, and a report relating the trade-offs to theory. Work split: J. Smouter – preprocessing, metric, brute force, BallTree; A. Nalla – TPS, MI-File, pivots; experiments, statistics and report are joint.

References

Amato, G., and Savino, P. (2008). Approximate similarity search in metric spaces using inverted files. In *Proc. 3rd International Conference on Scalable Information Systems (InfoScale)*.

Mohamed, H., and Marchand-Maillet, S. (2015). Scalable indexing for big data processing. In *Big Data: Algorithms, Analytics, and Applications*. Chapman and Hall/CRC.

Strack, B., DeShazo, J. P., Gennings, C., Olmo, J. L., Ventura, S., Cios, K. J., and Clore, J. N. (2014). Impact of HbA1c measurement on hospital readmission rates: analysis of 70,000 clinical database patient records. *BioMed Research International*, 2014, Article 781670. Data: UCI Machine Learning Repository, doi:10.24432/C5230J.